

Constrained Cache Design by a Council of Agents: A CHIA Loop That Measures and Gates Itself

Abstract—Choosing a cache hierarchy means setting the sets, ways, prefetcher, replacement policy and queue depth of every level at once, under a silicon budget and a power budget, with one simulation per design to judge it. We search 13 such knobs of a three-level hierarchy, 6.6 million designs, under area and power constraints of 4.0 mm² and 1.0 W, with a council of Gemini agents built as a CHIA loop: an analyst that diagnoses the current design, four specialists that each propose one move, a statistician with a Gaussian process, and a repair that shrinks a design that violated the power constraint. The council reaches the quality that random search attains after 1000 designs in 39 to 44 designs, 23 to 25 times fewer simulations. CHIA provided the nodes, the cluster and the profiler; we added what the loop needed and CHIA lacked: a Gemini node that bills thinking tokens, a ChampSim node that builds a whole hierarchy, an evaluation ledger, a spend view over the profiler log, and two gates in which a cheap calibrated decider judges which agent calls and which briefing sections are worth paying for. Applied, the gates reach the same final IPC for 22 to 30% less model spend. The blocks travel: on CHIA’s own CIRCT loop, the call gate cuts the assess stage from 16 to 13 agent turns while sending the same issues forward. Every table reproduces from the repository with one command and no API key.

I. INTRODUCTION

A cache hierarchy is usually set by sweeps and experience: one level at a time, one knob at a time, the rest held fixed. The knobs, however, do not act independently. A larger last level costs power that a smaller second level could have saved, a prefetcher that helps one level floods the next, and a constraint on silicon or power turns every gain into a trade-off. Since each design costs a simulation, the currency of the search is the number of designs measured.

CHIA [1] is a framework for building loops in which agents propose and tools measure, with every step a task on a cluster and a profiler recording it. We built such a loop for this problem: a council of agents that reads the simulator’s counters, diagnoses where the time and the power go, proposes designs and has them measured five at a time. Against random search over the same feasible space, the council reaches the quality that random search attains after 1000 designs in 39 to 44 designs, both under the constraints and without them.

Such loops are usually judged by that result alone, which says nothing about what the loop cost or where. The CHIA paper [1] names the gap: “data collection and profiling of both results and of the loop itself built in as first-class citizens”. We followed that principle throughout. The loop’s model calls, builds, simulations, candidates and gate decisions are all CHIA tasks, so one profiler log per run holds all of them, and the blocks we added either read that log or add events to it: a spend view that sums the log per node and model, a ledger that

scores every candidate the way search methods are compared, and two gates in which a cheap calibrated decider answers one yes-or-no question before an optional agent call or for each section of a briefing. Run in shadow, the gates only record; this showed that 57% of the specialists’ calls ended in a hold, and which of them a decider could have foreseen. Applied, the gates gave the same final IPC for 22 to 30% less model spend.

Sections II to VIII follow that order: the design problem, the council as a CHIA loop, the result against random search, the cost and the gates, the blocks on CHIA’s own CIRCT loop, what the instruments found, and the artifact.

II. THE DESIGN PROBLEM

The chip is ChampSim’s [2] stock core widened to 6-wide issue with a 384-entry reorder buffer, at 3.8 GHz in 22 nm, with DDR5-4800 over two channels. The search moves 13 knobs: the sets and ways of the L1D, L2 and LLC on per-level ladders, the prefetcher of each level, the replacement policy of the L2 and LLC, and the miss-handling depth of both, 6.6 million designs in all. Two constraints make this a design problem rather than a sweep: 4.0 mm² of cache silicon and 1.0 W of cache plus off-chip power, and a design may not run the workload slower than the stock chip does. The two constraints become known at different times, and that shapes the loop. Silicon and leakage come from CACTI 7 [3] at 22 nm and are known before a design is simulated, so a design that exceeds the area constraint is never run. Power is CACTI’s energies times the simulator’s hit and miss counters plus 15 pJ per bit moved off chip, so it is known only afterwards: a design that exceeds the power constraint is measured and read by the agents, but never becomes the incumbent. The stock chip itself draws 1.095 W on the workload and thus exceeds the power constraint, so the reference is the stock with its LLC halved: 1.96 mm², 0.805 W, IPC 0.935. The constraints are binding: 36% of the council’s candidates were refused after measurement, 21% by power and 19% by the speed floor, some by both. The area constraint rarely binds, since good designs are small; the best draws 0.67 W on 1.33 mm². A second regime drops both constraints.

The workload is the llama2 inference trace of the ChampSim ML suite, simulated at 1M warm-up and 2M instructions; 40 designs re-measured at 5M/25M rank the same (Spearman 0.986, top-5 overlap 80%). We chose the most demanding of six ML traces deliberately, the one with the most headroom over the stock and the deepest coupling between knobs, so that the council is tested where a sweep fails.




Fig. 1. One round of the council. Violet: the agents, tasks of the Gemini node; blue: the apparatus, CHIA’s ChampSim nodes and the ledger; green: the decider’s gate. Every box is a CHIA task.

III. THE COUNCIL AS A CHIA LOOP

One round. Fig. 1 shows one round. Round one measures the analyst’s opening design and 19 random feasible designs. Every later round measures one wave of five. A statistician, a Gaussian process over every design this run has measured, proposes its top pick; the analyst writes a sheet (a verdict per level with its evidence, the knobs that failed alone, the headroom under each constraint) and one whole design; four specialists, for prefetching, geometry, replacement and concurrency, each propose one move on their own knobs or hold. The wave takes at most five, in this order: the statistician’s pick, the analyst’s design, a gain refused last round by the power constraint and shrunk to fit, the specialists’ moves, and untried single-knob moves at the level the sheet blames, to fill what is left. The incumbent is the best feasible design measured, whoever proposed it, and a move already measured against it counts as a hold. One model serves every call, Gemini 3.1 Pro on Vertex at temperature 0.7 with a thinking budget of 4096 tokens.

What CHIA gave and what we added. Every box in Fig. 1 is a CHIA task on a Ray [4] cluster. Each run connects under its own name with its own profiler collector, so a run’s log is a complete record of it, and simulations run through CHIA’s own ChampSim run node. Three nodes did not exist, and we wrote them in CHIA’s layout. CHIA’s Vertex node drives a tool loop under a fixed generation config and leaves the thinking tokens out of its usage record, whereas our Gemini node makes one JSON call at a set temperature and thinking budget and records every token, the cost in dollars and the caller. CHIA’s build node writes one prefetcher at one level, whereas ours builds from a whole configuration. CHIA’s run node keeps a subset of the simulator’s record, less than the analyst reads, so a 7-line patch returns the whole of it, and the loop’s parse gives identical numbers. Every design measured is one ledger row and one `candidate` event in the log.

Protocol. The budget unit is a design: wall time at this fidelity shows no advantage either way, since the model’s minute per round offsets the fewer rounds. Random search draws 1000 feasible designs over three seeds; the council runs 60 designs, or 20 without a new incumbent, with four seeds per arm and regime. The target is the mean over random’s seeds of its best after 1000 designs; a seed that spends its budget under the target counts at the budget, 60 designs, with its final

TABLE I
THE COUNCIL AGAINST RANDOM SEARCH, MEANS OVER SEEDS. THE TARGET IS RANDOM’S MEAN BEST IPC AFTER 1000 DESIGNS; A COUNCIL SEED UNDER IT AT THE BUDGET COUNTS AS 60 DESIGNS. NEW BESTS: WHO PROPOSED THE NEW INCUMBENTS

	target	council designs	fewer	new bests found by		
	IPC	to the target	simulations	model	statistician	fills
constrained	1.4329	44	23 ×	62%	25%	13%
unconstrained	1.5001	39	25 ×	40%	36%	24%




Fig. 2. Constrained regime, means over seeds: the best feasible IPC as designs are measured, and the cumulative model spend of the two council arms beneath it. Random’s line is its 1000-design mean.

IPC beside it. Every arm is scored from the same ledger rows, random’s being rebuilt from its seed and the result table.

IV. AGAINST RANDOM SEARCH

The result. Random search (Table I, Fig. 2) climbs early, because a fifth of the feasible space beats the halved stock, and then stalls. The council passes random’s 1000-design quality at 44 designs under the constraints and at 39 without. Fifteen of its sixteen runs end above that quality; the remaining one ends at 1.497 against 1.5001.

Who found the improvements. The ledger records who proposed every new best design, and the answer differs between the two regimes. Under the constraints, the model produced most of them: its whole designs and its repairs of refused gains account for 62% of the new bests, the statistician for 25% and the single-knob fills for 13%. Unconstrained, the statistician led with 36%, the single-knob fills took 24%, and the model 40%.

Why. A constraint makes the objective discontinuous. A design can be faster and still be refused, and a surrogate fitted to IPC does not see that wall, whereas a model that reads the counters does, and can say which level to shrink. Without the constraints the landscape is smooth, and the surrogate’s interpolation is the better predictor. The constraints therefore decide who wins, not whether the search wins, and a ledger that records who won tells the two apart for each problem.

```

gate = CallGate(JevDecider(), threshold=0.4, mode="apply")
state = {"sheet": sheet_text(sheet)[:2500],
        "current_design": knobs_text(incumbent),
        "statistician_expects": expects[:800]}
questions = {}
for name in CONCERNS: # prefetch, geometry, replacement, ...
    questions[name] = {
        "instructions": f"Does the sheet give evidence "
                        f"for moving the {name} knobs?",
        "criteria": {"true": "its findings point here",
                    "false": "nothing does; a hold is free"}}
verdicts = gate_and_log(gate, state, questions, label="round")
for name in verdicts: # shadow: every call made
    if verdicts[name]["consult"]:
        ask_specialist(name) # apply: the rest not called

```

Listing 1. The call gate in front of the council’s specialists. The state and the questions belong to the loop, the rest to the block; ask_specialist is one task of the Gemini node.

Against Bayesian optimisation. We also ran a standard Bayesian optimisation, a random-forest surrogate with expected improvement and the same waves of five, two seeds per regime. It needs as many designs as the council to reach random’s 1000-design quality. Without the constraints it ends 1% above the council’s final IPC; under them it ends 1% below the council and 2.5% below the gated council, and half to two thirds of its designs came back in violation of the power constraint, against a third of the council’s. The same pattern holds: a surrogate predicts well on a smooth landscape, and where the constraint turns a faster design into an infeasible one, an agent that reads the counters does better.

V. CHEAPER AT THE SAME IPC

What the instruments showed. A run’s model bill is \$1.6 to \$1.7; what scales is the share saved, with more seeds, full fidelity, bigger loops, and every loop in which most optional calls end in a hold. The recipe above is the sixth version of the council, and the ledger determined most of what it contains: its rows showed that four rules of earlier versions never paid and that one setting stalled half the seeds. The per-call rows, priced by the spend block, showed what the bill was made of: input 49%, thinking 37%, answers 14%, and the analyst 44% of it. The specialists’ rows showed that 57% of their calls ended in a hold. Two levers therefore remained: the calls that end in a hold, and the briefing every call carries.

The gate. A gate in front of a call needs a verdict that is cheap, fast and calibrated, not a second agent. TypeSafe’s Jev [5] answers a typed yes-or-no question with a probability, at \$0.042 per million input tokens; a round of four questions costs about a tenth of a cent and returns within a second. The call gate asks, per specialist and round, whether the analyst’s sheet gives evidence for moving that concern’s knobs; Listing 1 is the whole of it. The context gate asks, per section of a briefing, whether the task needs it, judging from the caller’s one-line description of the section, and also rates the task’s difficulty, which sets the thinking budget. Both gates have two modes. In shadow, every call is still made and the verdict is recorded beside what the call then produced; in apply, the calls under the threshold are not made and the sections rated unneeded are dropped. The pair is the point: shadow prices the gate before apply collects the saving, so nothing is skipped on faith.

TABLE II

THE CALL GATE IN SHADOW OVER 304 SPECIALIST CALLS OF EIGHT RUNS (GAINS LOST: SKIPPED PROPOSALS THAT GAINED MORE THAN 0.005 IPC), THEN BOTH GATES APPLIED, 4 SEEDS PER ARM AND REGIME

<i>Shadow: what the gate would have skipped</i>				
threshold	skipped	holds	proposals	gains lost
0.3	93 (31%)	68	25	2
0.4	121 (40%)	92	29	2
0.5	147 (48%)	112	35	3
<i>Applied at 0.4, with the context gate; “all calls” is the ungated baseline</i>				
	constrained		unconstrained	
	all calls	gated	all calls	gated
model spend per run (\$)	1.63	1.27	1.71	1.20
model calls per run	47.2	34.0	49.8	31.0
final IPC, mean of seeds	1.512	1.537	1.521	1.513
less model spend	22%		30%	

What the shadow showed, and what applying it saved.

Table II: over the 304 specialist calls of the eight shadow runs, a threshold of 0.4 would have skipped 40%, 92 holds and 29 proposals, of which two gained more than 0.005 IPC; of the 20 gains above 0.005 among all 304 calls, 18 carried a verdict of 0.4 or more. The context gate drops 16 to 17% of the briefing tokens once shown a description per section; shown the sections’ first characters instead, it had rated the evidence unneeded, a negative result we report. Applied together, with a specialist’s thinking budget reduced when its task is not rated hard, the council consulted 2.4 to 2.8 specialists a round instead of four and reached the same final IPC for 22% less model spend under the constraints and 30% without (Fig. 2), the decider’s cost included. Nearly all of the saving, 94 to 98% of it, comes from the calls not made. The trimmed briefings cut a specialist call’s input by 7 to 12%, but the cost of a call made was flat, since thinking and answer lengths vary by more than that.

VI. THE BLOCKS TRAVEL

Everything the loop needed is a module and a test in CHIA’s tree (Table III), installed into the released package by one script; the tests pass inside it, and the CLI’s new format sums a log that CHIA’s own Vertex node wrote.

On CHIA’s own loop. CHIA’s CIRCT issue solver begins with an *assess* turn, in which an agent reads a GitHub issue and decides whether it is a bug with one clear fix; an issue that is not a bug is set aside after that turn, so it still costs the turn. We re-ran the turn under the blocks: the paper’s own prompt, the 16 issues of its Table 5, CHIA’s Vertex node under CHIA’s profiler, and the call gate asking the decider “does this issue report a defect?” before the turn, at 0.65. Table IV gives means over passes, since answers vary. The gate skipped the same three of the five non-bugs in both gated passes and none of the eleven bugs, so the paper’s seven issues proceed to reproduction as before, while turns and tokens fell 19% and the decider’s 16 decisions cost \$0.0008. The saving is bounded by the share of non-bugs in the stream, five of sixteen here; what the gate buys is that the expensive turn is spent only where it is needed, with no bug lost.

Six pull requests to CHIA. What the loop revealed about CHIA’s nodes went upstream as six pull requests, open at the

TABLE III
THE BLOCKS, WHAT EACH READS OR RECORDS, AND WHERE EACH WAS PROVEN

Block	What it reads or records	Proven on
Spend view and limit	Dollars and tokens by kind per node and model, from the events. <code>chia viz-profile -format spend;</code> <code>LLMSpend</code> stops a loop at a spend limit.	CHIA’s Vertex node under a local Ray; the CIRCT stage; the council’s own log
Evaluation ledger	One row per candidate: who, why, what it measured, which constraints it violated. From the rows: best-so-far, evaluations to a target, plateau, speedups, rank agreement.	Every arm of Sec. IV
Decider and call gate	One yes-or-no question per optional call, answered by Jev from the caller’s state; shadow records, apply skips.	The council (Sec. V); CIRCT: 16 to 13 turns
Context gate	Per briefing section, the decider judges the caller’s one-line description; it also rates the task’s difficulty, which sets the thinking budget.	The council: 16–17% fewer briefing tokens
Gemini JSON node	One JSON answer at a set temperature and thinking budget; the event carries every token kind, the cost in dollars and the caller.	Every model call of the council
ChampSim configuration build	Builds ChampSim from a whole configuration and returns CHIA’s own build result. A 7-line patch returns the simulator’s raw record.	Every build and run of the council; IPC equal to the loop’s own build

TABLE IV
CHIA’S ASSESS STAGE OVER THE PAPER’S 16 ISSUES (11 BUGS, 5 NOT), UNGATED AND GATED. AGREE: VERDICTS MATCHING THE PAPER’S, OF 16. TOKENS AND DOLLARS: THE SPEND VIEW OVER EACH PASS’S PROFILER LOG

	turns	agree	skipped		tokens	\$
			/16	bugs non-bugs		
ungated, 3 passes	16	14.3	0 of 11	0 of 5	37.7k	0.095
gated at 0.65, 2 passes	13	14.5	0 of 11	3 of 5	30.4k	0.081
change		−19%			−19%	−15%

time of writing. Four concern the Vertex node: a prompt that mentions an undeclared tool made it return success with an empty answer (14 of 16 issues came back blank); its usage record left out the thinking tokens; a rate limit raised at once, so a batch died on the first 429; and its generation settings could not be set per request. One resets the profiler when a collector starts and flushes it when it stops, so a driver that restarts the collector keeps its events. The sixth returns the ChampSim run node’s raw record, which had kept the first DRAM channel only and no merged misses.

VII. WHAT THE INSTRUMENTS FOUND

The model alone does not search. Without the statistician, the repairs and the fills, the specialists ran out of proposals at 25 to 28 designs, and the runs ended at IPC 1.15 to 1.17 in both regimes. **Grounding is a gauge, not a filter.** Asked whether the sheet supports each of 377 proposals, the decider’s verdict did not predict which of them gained, so it measures prompt discipline and should not be used to skip calls. **The decider cannot judge “one clear fix”** (10 of 11 CIRCT bugs were rated unclear), which is why the gate asks the narrower question.

VIII. ARTIFACT, RELATED WORK AND LIMITATIONS

Artifact. `python3 council_loop/demo.py` replays every table of this paper from the shipped records in un-

der a second, with no key, no cluster and no simulator; `council_loop/chia_blocks/check.sh` installs the released CHIA into a clean environment, copies the blocks in, applies the two patches and runs the 33 tests inside the package. The fixes are the six pull requests of Section VI; the blocks stay in this repository. The release tag `frozen-council` is the version the runs used. All of it is at https://github.com/mateobuscarons/chia_hackathon.

Related work. AgentDSE [6] and the CHIA studies [1] search architectural design spaces with reasoning agents and report the outcome. The ledger’s metrics are those with which predictive design-space exploration [7] and Bayesian optimisation [8] compare methods, evaluations to a target and best-so-far per evaluation, here scoring an agentic loop from its own records. FrugalGPT [9] lets a cheap model decide whether an expensive one is called; the call gate is that decision made by a calibrated decider [5] at the loop level, with a shadow mode that prices it before it is trusted.

Limitations. Four seeds per arm and regime bound the quality claim; the shadow counterfactual, 304 calls, is the larger sample. One trace at 1M/2M instructions carries the search, with a 40-design check at 5M/25M. The constraints are ours, taken from the stock chip’s own draw rather than from a product. Against the strongest random seed instead of the mean, the council needs 44 and 46 designs.

ACKNOWLEDGEMENTS

We thank Agustin Nahuel Coppari Hollmann (UC Berkeley) for support, the CHIA developers and the A³ hackathon organisers. Claude (Anthropic) assisted with the code and the drafting.

REFERENCES

- [1] A. Cui, F. Hermida-Rivera, J. Toubes, R. Gupta *et al.*, “CHIA: An open-source framework for principled, agentic AI-driven hardware/software co-design research,” *arXiv preprint arXiv:2606.27350*, 2026.
- [2] N. Gober, G. Chacon, L. Wang, P. V. Gratz, D. A. Jiménez, E. Teran, S. Pugsley, and J. Kim, “The Championship Simulator: Architectural simulation for education and competition,” *arXiv preprint arXiv:2210.14324*, 2022.
- [3] R. Balasubramonian, A. B. Kahng, N. Muralimanohar, A. Shafiee, and V. Srinivas, “CACTI 7: New tools for interconnect exploration in innovative off-chip memories,” *ACM Trans. Architecture and Code Optimization*, vol. 14, no. 2, pp. 14:1–14:25, 2017.
- [4] P. Moritz, R. Nishihara, S. Wang *et al.*, “Ray: A distributed framework for emerging AI applications,” in *Proc. 13th USENIX Symp. on Operating Systems Design and Implementation (OSDI)*, 2018, pp. 561–577.
- [5] TypeSafe AI, “System one documentation: the noul primitive, confidence and models,” <https://docs.typesafe.ai>, 2026.
- [6] C. Wang, J. C. Shi, D. Kong, D. S. Boning, Z. Wan, Y. Du, and V. J. Reddi, “AgentDSE: Reasoning-augmented architectural design space exploration,” *arXiv preprint arXiv:2606.21836*, 2026.
- [7] E. İpek, S. A. McKee, R. Caruana, B. R. de Supinski, and M. Schulz, “Efficiently exploring architectural design spaces via predictive modeling,” in *Proc. 12th Int. Conf. on Architectural Support for Programming Languages and Operating Systems (ASPLOS)*, 2006, pp. 195–206.
- [8] J. Snoek, H. Larochelle, and R. P. Adams, “Practical Bayesian optimization of machine learning algorithms,” in *Advances in Neural Information Processing Systems 25 (NIPS)*, 2012, pp. 2951–2959.
- [9] L. Chen, M. Zaharia, and J. Zou, “FrugalGPT: How to use large language models while reducing cost and improving performance,” *arXiv preprint arXiv:2305.05176*, 2023.